AI Ethical Principles Case Study

Assignment 1.3

Naga Venkata Satyanarayana Devarapalli

Indiana Wesleyan University | 2025

Developing a Responsible AI Framework: Ethical Principles for Enterprise AI Deployment

Fictional Company Overview: NovaMind AI Solutions

Company Profile

NovaMind AI Solutions is a mid-size technology company founded in 2019 and headquartered in Indianapolis, Indiana. The company specializes in developing AI-powered enterprise software products for organizations in the healthcare, financial services, and customer experience sectors.

NovaMind employs over 850 professionals, including a dedicated AI/ML engineering team of 150+ specialists focused on building intelligent systems that augment human decision-making. The company currently serves over 200 enterprise clients ranging from regional hospitals and credit unions to Fortune 500 financial institutions, generating approximately $120 million in annual revenue with 40% year-over-year growth.

Products & Services

Mission Statement

"To deliver intelligent, trustworthy, and human-centered AI products that empower organizations to make better decisions while upholding the highest ethical standards, regulatory compliance, and public trust."

Why NovaMind Needs an RAI Framework

As NovaMind's AI products directly impact critical domains—patient health, consumer financial access, and customer experiences—the company recognizes that ethical AI deployment is not merely a compliance requirement but a business imperative. Clients increasingly demand demonstrable AI ethics practices, regulators are tightening oversight (EU AI Act, FDA AI/ML guidance), and public trust in AI remains fragile. A comprehensive Responsible AI framework enables NovaMind to:

- AI-powered clinical decision support software for healthcare providers
- Machine learning-based credit risk assessment platforms for financial institutions
- Intelligent conversational AI solutions for enterprise customer service
- Internal AI productivity tools for document analysis and knowledge management

- Mitigate legal and reputational risks from biased or opaque AI decisions
- Build competitive advantage through demonstrable trustworthiness
- Meet regulatory requirements across multiple jurisdictions proactively
- Attract and retain talent who value ethical technology development
- Maintain long-term client relationships built on accountability

Compliance & Certifications

HIPAA compliant, SOC 2 Type II certified, GDPR-ready infrastructure, ISO 27001 in progress. NovaMind maintains a dedicated AI Ethics Board comprising internal leaders, external academics, and industry advisors.

AI Use Cases Under Consideration (2–4 Specific Use Cases)

Use Case	Description	Target Users	AI Techniques	Ethical Considerations
1. Clinical Decision Support System (CDSS)	An AI-powered diagnostic assistant that analyzes electronic health records (EHRs), lab results, radiology images, and patient history to recommend potential diagnoses and suggest evidence-based treatment options. The system provides confidence scores and highlights contributing factors for each recommendation.	Physicians, nurses, radiologists at partner hospitals and clinics	Deep learning (CNNs for imaging), gradient-boosted decision trees for risk prediction, NLP for clinical notes extraction	Patient safety, informed consent, algorithmic bias across demographics, physician over-reliance on AI recommendations, data privacy (HIPAA), liability for misdiagnosis

Use Case	Description	Target Users	AI Techniques	Ethical Considerations
2. Credit Risk Scoring Engine	A machine learning model that evaluates loan and credit card applicants' creditworthiness by analyzing traditional financial indicators (credit history, debt-to-income ratio) alongside alternative data sources (utility payments, rent history) to provide more inclusive credit assessments while managing institutional risk.	Loan officers, underwriters, compliance teams at banking and credit union clients	Ensemble methods (XGBoost, Random Forest), logistic regression for baseline, feature importance analysis	Fair lending compliance, disparate impact on protected classes, transparency in adverse action notices, data quality for alternative sources, potential to perpetuate historical bias
3. Intelligent Customer Service Chatbot	A conversational AI system built on large language model (LLM) technology that handles tier-1 customer inquiries across multiple channels (web, mobile, email). The	End consumers interacting with NovaMind's enterprise clients; customer service managers overseeing	Large language models (fine-tuned), intent classification, sentiment analysis, retrieval-augmented generation (RAG)	AI disclosure (users knowing they're speaking with AI), misinformation risk from hallucinations, accessibility for diverse users, emotional manipulation

Use Case	Description	Target Users	AI Techniques	Ethical Considerations
	chatbot resolves common questions, processes routine requests, and intelligently escalates complex issues to human agents with full conversation context.	bot performance		concerns, data retention of conversations
4. Internal Document Summarization Tool	An internal productivity tool that automatically generates concise summaries of lengthy reports, meeting transcripts, research papers, and policy documents. Employees can request summaries at varying detail levels to accelerate information consumption and knowledge	Internal NovaMind employees across all departments	Transformer-based summarization, extractive and abstractive approaches, keyword extraction	Information accuracy, potential loss of nuance in summaries, over-reliance on AI-generated content, intellectual property considerations for summarized external research

Use Case	Description	Target Users	AI Techniques	Ethical Considerations
	sharing across the organization.			

NovaMind's AI Ethical Principles

Based on research into established AI ethics frameworks from leading organizations (Google AI Principles, Microsoft Responsible AI Standard, IBM's Everyday Ethics for AI, OECD AI Principles, and the IEEE Ethically Aligned Design), NovaMind has adopted the following six core ethical principles to guide all AI development and deployment:

1. Fairness & Non-Discrimination

Principle: NovaMind AI systems must not create or reinforce unfair bias against individuals or groups based on protected characteristics including race, gender, age, disability, religion, or socioeconomic status.

Application: All models undergo bias testing across demographic subgroups before deployment. The Credit Risk Scoring Engine specifically requires disparate impact analysis and adverse action explainability to comply with ECOA and FCRA. The Clinical Decision Support System must demonstrate equitable performance across patient demographics to prevent diagnostic disparities.

2. Transparency & Explainability

Principle: NovaMind AI systems must be transparent about their capabilities, limitations, and decision-making processes. Users should receive explanations appropriate to their context and comprehension level.

Application: High-risk systems (CDSS, Credit Risk) require full explainability with methods like SHAP values, feature importance, and natural language explanations. Customer-facing AI must clearly disclose its non-human nature. All systems are documented with Model Cards describing architecture, performance, limitations, and intended use.

3. Privacy & Data Governance

Principle: NovaMind respects individual privacy rights and processes personal data only with appropriate authorization, purpose limitation, and security safeguards aligned with HIPAA, GDPR, CCPA, and applicable regulations.

Application: Patient data in CDSS is accessed only under covered entity agreements with minimum necessary access. Credit data processing follows FCRA requirements with consumer consent. Customer chatbot conversations are retained only as long as necessary with clear disclosure. All training data undergoes privacy impact assessments.

4. Safety & Reliability

Principle: NovaMind AI systems must operate safely, reliably, and as intended. Systems must have appropriate safeguards to prevent harm, including human oversight mechanisms for high-stakes decisions.

Application: The CDSS operates as a decision-support tool only—physicians retain full decision authority. Credit scoring includes manual review thresholds for borderline cases. The chatbot has escalation triggers for sensitive topics (medical emergencies, financial hardship). All systems undergo adversarial testing and have failsafe mechanisms.

5. Accountability & Governance

Principle: Clear ownership and responsibility structures exist for all AI systems. NovaMind maintains mechanisms for addressing AI-related concerns, errors, and harms through its AI Ethics Board and incident response procedures.

Application: Each AI product has a designated Responsible AI Owner. The AI Ethics Board reviews all high-risk systems before deployment and quarterly thereafter. An AI incident response plan defines escalation procedures. External stakeholders can submit concerns through a dedicated AI ethics contact channel.

6. Human Autonomy & Oversight

Principle: NovaMind AI systems augment rather than replace human judgment. Affected individuals have the right to human review of significant AI-driven decisions and the ability to contest outcomes they believe are incorrect or unfair.

Application: CDSS recommendations require physician approval before implementation. Credit applicants denied based on AI scoring can request human review. Chatbot users can always request transfer to a human agent. No fully autonomous decision-making is permitted for high-risk use cases.

Mapping Ethical Principles to AI Use Cases

The following chart demonstrates how NovaMind's ethical principles apply differently across use cases based on risk level, stakeholder impact, and regulatory context:

Ethical Principle	Clinical Decision Support	Credit Risk Scoring	Customer Service Chatbot	Document Summarization
Fairness	Critical – Must perform equitably across racial, age, and gender groups to prevent diagnostic bias	Critical – Legal requirement under fair lending laws; disparate impact testing mandatory	Moderate – Should provide equitable service quality regardless of user demographics or language	Low – Internal tool with minimal fairness risk; ensure accessible to all employees
Transparency	Critical – Physicians must understand diagnostic reasoning; patients may request explanations	Critical – Legal requirement to explain adverse credit decisions; model must be interpretable	Moderate – Must disclose AI nature; helpful to explain capabilities and limitations	Low – Nice-to-have; employees can verify against source documents
Privacy	Critical – Protected health information (PHI) under HIPAA;	High – Consumer financial data under	Moderate – Conversation data retained with	Low – Internal documents only; standard corporate data

Ethical Principle	Clinical Decision Support	Credit Risk Scoring	Customer Service Chatbot	Document Summarization
	strict access controls required	FCRA/GLBA; clear consent and purpose limitation needed	disclosure; minimize personal data collection	governance applies
Safety	Critical – Incorrect recommendations could harm patients; robust validation and human oversight essential	High – Errors could unfairly deny financial access; safeguards against systematic failures needed	Moderate – Incorrect information could frustrate users; escalation for sensitive topics required	Low – Errors easily caught by comparing to source; minimal safety risk
Accountability	Critical – Clear liability chain; physician remains accountable for medical decisions	Critical – Institution remains legally responsible for lending decisions; audit trail required	Moderate – Clear ownership of bot behavior; escalation procedures documented	Low – Internal productivity tool; standard product ownership

Ethical Principle	Clinical Decision Support	Credit Risk Scoring	Customer Service Chatbot	Document Summarization
Human Oversight	Critical – AI recommends only; physician decides. No autonomous medical decisions	High – Human review for borderline cases and appeals; no fully automated denials for large loans	Moderate – Human escalation always available; complex issues routed to agents	Low – Employees exercise judgment about summary accuracy before acting on content

Research: How Major Technology Companies Approach AI Ethics

NovaMind's ethical framework was informed by studying how leading technology companies have established and communicated their AI ethics approaches:

Google – AI Principles (2018)

Google published seven AI principles emphasizing: being socially beneficial, avoiding unfair bias, being built and tested for safety, being accountable to people, incorporating privacy design principles, upholding high standards of scientific excellence, and being made available for uses that accord with these principles. Google also lists AI applications they will not pursue (weapons, surveillance violating norms, technologies causing harm). Their approach is notable for its public commitment to not weaponize AI and for establishing an Advanced Technology External Advisory Council (though it was dissolved after controversy).

Microsoft – Responsible AI Standard (2022)

IBM – Everyday Ethics for AI (2019)

IBM's framework focuses on five focal areas: accountability, value alignment, explainability, fairness, and user data rights. Their approach is distinctive for its emphasis on practical design guidance—providing AI designers and developers with concrete questions to ask at each stage of development. IBM also created open-source toolkits (AI Fairness 360, AI Explainability 360, Adversarial Robustness Toolbox) that operationalize ethical principles into testable code. Their "Ethics by Design" approach embeds ethical considerations into the development workflow rather than treating them as post-hoc compliance checks.

Microsoft's framework is highly operationalized with specific requirements organized around six principles: fairness, reliability & safety, privacy & security, inclusiveness, transparency, and accountability. Their Responsible AI Standard v2 provides concrete implementation requirements for product teams including mandatory impact assessments, fairness testing, and transparency documentation. Microsoft also developed tools like Fairlearn and InterpretML to support implementation. Their approach stands out for translating principles into specific, auditable product requirements.

Meta – Responsible AI Practices

Meta's approach centers on five pillars: privacy, fairness, robustness, transparency, and accountability. They are notable for publishing research on responsible AI (including bias detection in large language models) and for creating system cards that explain how AI systems work to users. Meta's cross-functional Responsible AI team works with product teams to integrate ethical considerations. Their transparency reports detail AI content moderation decisions and system performance metrics. However, Meta has also faced significant criticism regarding algorithmic amplification and content recommendation ethics.

Key Takeaway for NovaMind

All four companies share common principles (fairness, transparency, accountability, safety) but differ in operationalization. NovaMind adopts Microsoft's approach to creating specific, auditable requirements while incorporating IBM's emphasis on practical tooling and Google's clarity on prohibited applications.

RAI Framework Implementation Roadmap

NovaMind's responsible AI framework will be implemented in phases to ensure thorough integration across all AI products and organizational processes:

Phase	Timeline	Key Activities	Success Metrics
Phase 1: Foundation	Q1 2025	Establish AI Ethics Board; publish internal ethical principles; develop risk classification framework; create AI Impact Assessment template; conduct baseline ethics audit of existing products	Ethics Board operational; principles ratified; all existing products classified by risk tier
Phase 2: Tooling & Process	Q2 2025	Implement bias testing pipeline; deploy explainability tools (SHAP, LIME integration); develop Model Card and Data Card templates; establish pre-deployment review gates; train all AI/ML engineers on RAI practices	100% of new models go through bias testing; explainability integrated into CI/CD pipeline; 100% engineer training completion
Phase 3: Deployment	Q3 2025	Apply framework to all new AI products; retrofit high-risk existing products; implement user-facing transparency features (AI disclosure, explanation interfaces); establish external feedback channels	All high-risk products fully compliant; transparency notices deployed for customer-facing AI; feedback channel operational

Phase	Timeline	Key Activities	Success Metrics
Phase 4: Maturity	Q4 2025+	Conduct first annual third-party RAI audit; publish transparency report; refine framework based on learnings; expand to cover new AI modalities (generative AI, multimodal); engage in industry standards development	Third-party audit completed; public transparency report published; framework updated based on findings

---

Conclusion: Building Ethical AI at NovaMind

Summary of Key Decisions

- NovaMind adopted six core ethical principles (Fairness, Transparency, Privacy, Safety, Accountability, Human Oversight) informed by industry leaders' frameworks
- Ethical requirements are proportional to risk—high-risk systems (healthcare, finance) face the most rigorous requirements while internal tools follow baseline standards
- The company established an AI Ethics Board with cross-functional and external membership for independent oversight
- Implementation follows a phased approach to ensure sustainable adoption without disrupting existing operations
- Open-source tools and industry-standard documentation practices (Model Cards, Data Cards) will support consistent implementation

Personal Reflections

Developing this case study highlighted that AI ethics is not a one-time checklist but a continuous organizational commitment. The most challenging aspect was determining how to balance competing principles—for example, when maximum transparency might conflict with privacy, or when safety requirements (human oversight) might reduce the efficiency benefits that motivated AI adoption in the first place.

I also found that studying real company frameworks reveals a gap between published principles and actual practice. Companies that have faced public controversy (Meta's algorithmic amplification, IBM discontinuing facial recognition) demonstrate that principles alone are insufficient without robust enforcement mechanisms and willingness to make difficult business decisions when ethics require it.

Going forward, the key challenge for NovaMind—and for any AI company—is maintaining ethical rigor as competitive pressures and rapid technological advancement (particularly generative AI) push toward faster deployment with less oversight.

References: Google AI Principles (2018); Microsoft Responsible AI Standard v2 (2022); IBM Everyday Ethics for AI (2019); Meta Responsible AI Practices; OECD Principles on AI (2019); IEEE Ethically Aligned Design (2019); EU AI Act (2024); NIST AI Risk Management Framework (2023)