

Responsible AI Framework: Explainability & Transparency

Naga Venkata Satyanarayana Devarapalli

Indiana Wesleyan University | 2025

Company Overview: NovaMind AI Solutions

About NovaMind

NovaMind AI Solutions is a mid-size technology company specializing in AI-powered enterprise software for healthcare, finance, and customer service industries. Founded in 2019 and headquartered in Indianapolis, IN, NovaMind serves over 200 enterprise clients across the United States.

Mission

To deliver intelligent, trustworthy, and human-centered AI products that empower organizations to make better decisions while maintaining ethical standards and regulatory compliance.

Key Facts

- **Employees:** 850+
- **Annual Revenue:** \$120M
- **AI/ML Engineers:** 150+
- **Industries Served:** Healthcare, Finance, Customer Service
- **Compliance:** HIPAA, SOC 2, GDPR-ready

AI Use Cases Under Consideration

1. Clinical Decision Support

An AI system that assists physicians by analyzing patient records, lab results, and medical imaging to recommend potential diagnoses and treatment plans. Used in partner hospitals and clinics.

2. Credit Risk Scoring

A machine learning model that evaluates loan applicants' creditworthiness by analyzing financial history, income patterns, and alternative data sources for banking clients.

3. Customer Service Chatbot

A natural language processing (NLP) chatbot that handles tier-1 customer inquiries, routes complex issues to human agents, and provides 24/7 support for enterprise clients.

4. Internal Document Summarization

An internal AI tool that automatically summarizes lengthy reports, meeting notes, and research papers for employees to improve productivity and knowledge sharing.

Explainability Requirements by Use Case


AI Use Case	Explainability Required?	Level	Rationale
Clinical Decision Support	Yes – Mandatory	High	Directly impacts patient health outcomes. Physicians need to understand AI reasoning to make informed medical decisions. Regulatory requirements (FDA, HIPAA) mandate transparency in clinical tools. Errors can result in patient harm or legal liability.
Credit Risk Scoring	Yes – Mandatory	High	Subject to fair lending regulations (ECOA, FCRA). Applicants have a legal right to understand why credit was denied. Must demonstrate absence of discriminatory bias. Regulators require model interpretability for audit purposes.
Customer Service Chatbot	Partial – Recommended	Medium	Lower risk since humans handle escalations. Transparency helps users understand bot limitations. Useful for debugging incorrect responses. Not legally required but improves trust and user satisfaction.
Internal Document Summarization	No – Optional	Low	Used internally for productivity; no external stakeholder impact. Employees can verify summaries

AI Use Case	Explainability Required?	Level	Rationale
			against source documents. Low risk of harm from errors. Explainability is nice-to-have but not critical for deployment.

Key Principle: Explainability requirements scale with the potential impact on individuals' rights, health, financial status, and the degree of regulatory oversight applicable to the domain.



Process for Evaluating Transparency & Explainability



Step	Description	Key Activities
1. Risk Classification	Categorize the AI system by impact level (High/Medium/Low)	Assess potential harm, regulatory scope, affected populations, reversibility of decisions
2. Stakeholder Identification	Identify who needs explanations and at what level of detail	Map end users, regulators, internal teams, affected individuals; determine technical vs. non-technical needs
3. Explainability Method Selection	Choose appropriate explanation techniques based on model type and audience	Select from SHAP, LIME, feature importance, decision trees, counterfactual explanations, natural language summaries
4. Testing & Validation	Verify that explanations are accurate, consistent, and comprehensible	User testing with target audiences, accuracy verification of explanations, consistency checks across similar inputs

Step	Description	Key Activities
5. Ongoing Monitoring	Continuously evaluate explainability as models evolve	Regular audits, drift detection, user feedback collection, re-validation after model updates



Documentation to Promote Transparency

AI Development Stage	Documentation Artifact	Purpose & Contents
Problem Definition & Planning	AI Impact Assessment	Documents intended use, affected populations, potential risks, explainability requirements, and ethical considerations before development begins.
Data Collection & Preparation	Data Card / Datasheet	Records data sources, collection methods, known biases, representativeness, preprocessing steps, and limitations of training data.
Model Development	Model Card	Describes model architecture, training procedures, hyperparameters, performance metrics, intended use cases, known limitations, and fairness evaluations.
Testing & Validation	Evaluation Report	Documents test results, bias audits, explainability test outcomes, edge case analysis, and comparison against baseline approaches.
Deployment	Transparency Notice / User Guide	Provides end users with clear information about AI involvement, how decisions are made, how to request explanations, and how to appeal automated decisions.

AI Development Stage	Documentation Artifact	Purpose & Contents
Post-Deployment Monitoring	Audit Log & Monitoring Report	Tracks model performance over time, records drift incidents, documents retraining events, logs user feedback, and maintains compliance records.

Guiding Standard: Documentation follows frameworks inspired by Google's Model Cards, Microsoft's Datasheets for Datasets, and the EU AI Act's transparency obligations.



Concluding Reflections: Difficulties & Questions

Challenges Encountered

- **Balancing complexity and clarity:** Determining the right level of explanation for different audiences is genuinely difficult. A data scientist and a patient need very different types of explanations from the same system.
- **Explainability vs. performance trade-off:** More interpretable models (e.g., decision trees) often sacrifice predictive accuracy compared to black-box models (e.g., deep neural networks). Deciding where this trade-off is acceptable requires careful ethical judgment.
- **Defining "sufficient" transparency:** There is no universal standard for what makes an AI system "transparent enough." Different regulators, industries, and

Open Questions

- How do we measure whether an explanation is truly understood by its intended audience, rather than just provided?
- As AI models become more complex (e.g., large language models), will current explainability techniques remain adequate?
- How should companies handle situations where full transparency might expose proprietary algorithms or create security vulnerabilities?
- What role should industry self-regulation play versus government-mandated transparency requirements?

stakeholders have varying expectations.

These reflections represent my personal observations from engaging with the course material on AI explainability and transparency. The intersection of technical capability and ethical responsibility remains one of the most challenging and important aspects of responsible AI development.

© 2025 Naga Venkata Satyanarayana Devarapalli | RAI Framework – Explainability & Transparency

Indiana Wesleyan University | Responsible AI Workshop