Responsible AI Framework: Explainability & Transparency

Naga Venkata Satyanarayana Devarapalli

Indiana Wesleyan University | Assignment 2.3 | 2025

Incorporating explainability and transparency considerations into a comprehensive Responsible AI (RAI) framework for enterprise AI deployment.

Slide 1: Company Overview – NovaMind AI Solutions

Company Profile

NovaMind AI Solutions is a mid-size technology company founded in 2019 and headquartered in Indianapolis, Indiana. The company specializes in developing AI-powered enterprise software products for organizations in the healthcare, financial services, and customer experience sectors. NovaMind employs over 850 professionals, including a dedicated AI/ML engineering team of 150+ specialists focused on building intelligent systems that augment human decision-making.

Industry Presence

NovaMind currently serves over 200 enterprise clients ranging from regional hospitals and credit unions to Fortune 500 financial institutions. The company generates approximately $120 million in annual revenue and has experienced 40% year-over-year

Mission Statement

"To deliver intelligent, trustworthy, and human-centered AI products that empower organizations to make better decisions while upholding the highest ethical standards, regulatory compliance, and public trust."

Core Values Relevant to RAI

- **Transparency:** All AI products must be understandable to their intended users
- **Accountability:** Clear ownership of AI decisions at every organizational level
- **Fairness:** Commitment to identifying and mitigating bias in all AI systems
- **Privacy:** Strict data governance aligned with HIPAA, GDPR, and CCPA
- **Human-Centricity:** AI augments human expertise

growth due to increasing enterprise demand for AI-driven automation and analytics.

rather than replacing human judgment

Compliance & Certifications

HIPAA compliant, SOC 2 Type II certified, GDPR-ready infrastructure, ISO 27001 in progress. The company maintains a dedicated AI Ethics Board comprising internal leaders, external academics, and industry advisors.

Slide 2: AI Use Cases Under Consideration (2–4 Specific Use Cases)

Use Case	Description	Target Users	AI Techniques	Deployment Context
1. Clinical Decision Support System (CDSS)	An AI-powered diagnostic assistant that analyzes electronic health records (EHRs), lab results, radiology images, and patient history to recommend potential diagnoses and suggest evidence-based treatment options. The system provides confidence scores and highlights contributing factors for each recommendation.	Physicians, nurses, radiologists at partner hospitals and clinics	Deep learning (CNNs for imaging), gradient-boosted decision trees for risk prediction, NLP for clinical notes extraction	External – deployed in 15 partner healthcare facilities with direct patient impact

Use Case	Description	Target Users	AI Techniques	Deployment Context
2. Credit Risk Scoring Engine	A machine learning model that evaluates loan and credit card applicants' creditworthiness by analyzing traditional financial indicators (credit history, debt-to-income ratio) alongside alternative data sources (utility payments, rent history) to provide more inclusive credit assessments while managing institutional risk.	Loan officers, underwriters, compliance teams at banking and credit union clients	Ensemble methods (XGBoost, Random Forest), logistic regression for baseline comparison, feature importance analysis	External – integrated into client banking platforms affecting consumer credit decisions
3. Intelligent Customer Service Chatbot	A conversational AI system built on large language model (LLM) technology that handles tier-1 customer inquiries across multiple channels (web, mobile, email). The	End consumers interacting with NovaMind's enterprise clients; customer service managers overseeing	Large language models (fine-tuned), intent classification, sentiment analysis, retrieval-augmented generation (RAG)	External – customer-facing deployment handling 50,000+ interactions daily across 30 enterprise clients

Use Case	Description	Target Users	AI Techniques	Deployment Context
	chatbot resolves common questions, processes routine requests, and intelligently escalates complex issues to appropriate human agents with full conversation context.	bot performance		
4. Internal Document Summarization Tool	An internal productivity tool that automatically generates concise summaries of lengthy reports, meeting transcripts, research papers, and policy documents. Employees can request summaries at varying detail levels (executive brief, detailed summary, key action items) to	Internal NovaMind employees across all departments (engineering, legal, sales, executives)	Transformer-based summarization models, extractive and abstractive summarization, keyword extraction	Internal only – no external stakeholder exposure; employees verify outputs against source documents

Use Case	Description	Target Users	AI Techniques	Deployment Context
	accelerate information consumption and knowledge sharing across the organization.			

Slide 3: Chart – Which AI Use Cases Require Explainability

The following chart classifies each AI use case by whether explainability is required, based on regulatory obligations, risk of harm, stakeholder needs, and industry best practices. This assessment draws on research into how companies like Google, Microsoft, IBM, and Meta approach AI transparency in similar domains.

AI Use Case	Explainability Required?	Risk Level	Regulatory Drivers	Rationale
Clinical Decision Support	Yes – Mandatory	High	FDA guidance on Clinical Decision Support Software, HIPAA, state medical practice regulations	This system directly impacts patient health and safety. Physicians must understand why the AI recommends a specific diagnosis to exercise professional medical judgment. The FDA requires that CDS tools provide sufficient information for clinicians to independently evaluate recommendations. Without explainability, physicians cannot verify AI outputs,

AI Use Case	Explainability Required?	Risk Level	Regulatory Drivers	Rationale
				creating unacceptable liability and patient safety risks. Google Health and IBM Watson Health both emphasize full clinical transparency in their healthcare AI products.
Credit Risk Scoring	Yes – Mandatory	High	Equal Credit Opportunity Act (ECOA), Fair Credit Reporting Act (FCRA), EU AI Act (high-risk classification), Consumer Financial Protection Bureau (CFPB) guidance	Federal law requires lenders to provide specific reasons when credit is denied or terms are adverse. The model must produce interpretable outputs explaining which factors contributed to a score. This is not optional—it is a legal requirement. Additionally, the system must be auditable for discriminatory bias against protected classes. Microsoft's

AI Use Case	Explainability Required?	Risk Level	Regulatory Drivers	Rationale
				Responsible AI Standard and IBM's AI Fairness 360 toolkit both provide frameworks for ensuring financial AI explainability.
Customer Service Chatbot	Partial – Recommended	Medium	FTC guidance on AI disclosure, EU AI Act transparency requirements for AI systems interacting with humans	While individual chatbot responses carry lower risk than medical or financial decisions, transparency is important for several reasons: (1) Users should know they are interacting with AI, not a human (required by EU AI Act Article 52); (2) When the bot provides incorrect information, transparent logging helps identify and fix issues; (3) Escalation logic should be explainable to service managers. Google's AI

AI Use Case	Explainability Required?	Risk Level	Regulatory Drivers	Rationale
				Principles and Meta's approach to conversational AI both emphasize user awareness and bot disclosure.
Internal Document Summarization	No – Optional	Low	No specific regulatory requirements for internal productivity tools	This tool is used exclusively by internal employees for convenience. Users can always compare summaries against original source documents to verify accuracy. The consequences of an imperfect summary are minimal (re-reading the source). No external stakeholders are affected, and no legally binding decisions are made based solely on these summaries. However, NovaMind still maintains basic documentation of

AI Use Case	Explainability Required?	Risk Level	Regulatory Drivers	Rationale
				the model's capabilities and limitations as an internal best practice, consistent with Microsoft's approach to internal AI tool governance.

Research References on Industry Transparency Approaches: Google AI Principles (2018) emphasize that AI systems should "be accountable to people" and provide "appropriate transparency and explanations." Microsoft's Responsible AI Standard (2022) requires transparency documentation for all AI systems proportional to risk. IBM's Everyday Ethics for AI framework categorizes explainability needs by use-case impact. Meta's Responsible AI team publishes system cards detailing how AI systems work for user-facing products.

Slide 4: Chart – Process for Evaluating AI Transparency & Explainability

NovaMind follows a structured 5-step evaluation process to determine whether each AI product meets the company's transparency and explainability standards. This process is applied at the design phase, before deployment, and periodically post-deployment.

Step 1: Risk & Impact Classification → Step 2: Stakeholder Mapping → Step 3: Explainability Method Selection → Step 4: Comprehension Testing → Step 5: Continuous Monitoring

Evaluation Step	Detailed Description	Key Activities & Deliverables	Responsible Team
Step 1: Risk & Impact Classification	Every AI product is categorized into a risk tier (High / Medium / Low) based on its potential to cause harm, affect rights, or trigger regulatory obligations. This assessment considers the severity of potential errors, the vulnerability of affected populations, the reversibility of AI-driven	Complete risk assessment questionnaire; map regulatory requirements; classify affected populations by vulnerability; produce Risk Tier Assignment document reviewed by AI Ethics Board	Product Management + Legal + AI Ethics Board

Evaluation Step	Detailed Description	Key Activities & Deliverables	Responsible Team
	decisions, and applicable legal frameworks.		
Step 2: Stakeholder Mapping	Identify all parties who need explanations about the AI system's behavior, and determine what type of explanation each stakeholder group requires. Different audiences need different explanation styles—a data scientist needs technical model details, while an end user needs a plain-language summary of how a decision was reached.	Create stakeholder registry; categorize explanation needs (technical vs. non-technical vs. regulatory); define explanation format requirements (visual, textual, interactive); document expected user literacy levels	Product Management + UX Research + Compliance
Step 3: Explainability Method Selection	Based on the model type, risk tier, and stakeholder needs, select appropriate explainability techniques. For high-risk systems, multiple complementary methods may be required. The selection must balance explanation fidelity (how accurately the explanation reflects the model's actual reasoning) with comprehensibility (how	Evaluate applicable XAI methods (SHAP values, LIME, feature importance rankings, attention visualization, counterfactual explanations, decision tree surrogates, natural language explanation generation); prototype explanations;	AI/ML Engineering + Data Science + UX Design

Evaluation Step	Detailed Description	Key Activities & Deliverables	Responsible Team
	easily the target audience can understand it).	document method selection rationale	
Step 4: Comprehension Testing	Before deployment, verify that the chosen explanations are actually understood by their intended audiences. An explanation that is technically accurate but incomprehensible to its target user fails the transparency goal. Testing includes both accuracy validation (does the explanation faithfully represent the model?) and user comprehension studies.	Conduct user comprehension studies with representative stakeholders; verify explanation accuracy against model internals; test edge cases where explanations might be misleading; measure user trust calibration; iterate on explanation design based on feedback	UX Research + QA + External User Panels
Step 5: Continuous Monitoring	After deployment, continuously monitor whether explanations remain accurate and relevant as models are updated, data distributions shift, and user needs evolve. Schedule periodic re-evaluations and establish triggers for immediate reassessment (e.g., model retraining, regulatory changes, user	Quarterly transparency audits; automated drift detection alerts; user feedback collection and analysis; re-validation after every model update; annual third-party explainability audit for high-risk systems; maintain explanation	AI Operations + Compliance + AI Ethics Board

Evaluation Step	Detailed Description	Key Activities & Deliverables	Responsible Team
	complaints about AI behavior).	versioning aligned with model versions	

Industry Alignment: This evaluation framework is informed by Microsoft's Responsible AI Impact Assessment process, Google's PAIR (People + AI Research) guidelines for human-centered AI design, and the NIST AI Risk Management Framework (AI RMF 1.0, 2023) which recommends proportionate transparency measures based on context and risk.

Slide 5: Chart – Documentation During AI Development to Promote Transparency

NovaMind requires specific documentation artifacts at each stage of AI development. This documentation creates an auditable record of decisions, promotes organizational accountability, and ensures that transparency is built into the development lifecycle rather than added as an afterthought.

AI Development Stage	Required Documentation	Detailed Contents & Purpose	Audience
1. Problem Definition & Planning	AI Impact Assessment (AIA)	Defines the business problem, documents why AI is the appropriate solution approach, identifies all affected stakeholders, assesses potential harms and benefits, establishes explainability requirements based on risk classification, defines success metrics and acceptable error thresholds, and outlines the human oversight model. Must be approved by the AI Ethics Board before development begins.	Product owners, AI Ethics Board, legal/compliance, executive leadership

AI Development Stage	Required Documentation	Detailed Contents & Purpose	Audience
2. Data Collection & Preparation	Data Card / Datasheet for Datasets	Comprehensively records all data sources used (including third-party data), collection methodologies, temporal coverage, demographic representation analysis, known biases and gaps, data cleaning and preprocessing steps applied, feature engineering decisions and rationale, privacy protections implemented (anonymization, encryption), and data retention/deletion policies. Based on Gebru et al.'s "Datasheets for Datasets" framework adopted by companies like Google and Microsoft.	Data engineers, ML engineers, compliance auditors, privacy officers
3. Model Development & Training	Model Card	Documents model architecture choices and alternatives considered, training procedures and hyperparameter selections, performance metrics across different demographic subgroups (disaggregated evaluation), intended use cases and	ML engineers, peer reviewers, AI Ethics Board, external auditors

AI Development Stage	Required Documentation	Detailed Contents & Purpose	Audience
		explicitly out-of-scope uses, known limitations and failure modes, fairness evaluations and bias testing results, environmental impact (compute resources used), and version history. Follows the Model Cards framework introduced by Mitchell et al. (2019) and adopted by Google, Hugging Face, and others.	
4. Testing & Validation	Evaluation Report & Bias Audit	Documents comprehensive testing results including accuracy on holdout data, performance across demographic subgroups (fairness metrics: equalized odds, demographic parity, etc.), explainability testing outcomes (do explanations accurately reflect model behavior?), adversarial robustness testing, edge case and stress test results, comparison against baseline/incumbent approaches, red team findings, and sign-off criteria for deployment readiness.	QA teams, ML engineers, AI Ethics Board, product management, external auditors for high-risk systems

AI Development Stage	Required Documentation	Detailed Contents & Purpose	Audience
5. Deployment & Release	Transparency Notice & User-Facing Documentation	Provides end users with clear, accessible information including: notification that AI is being used in the interaction/decision, plain-language description of what factors the AI considers, how to access more detailed explanations of individual decisions, how to contest or appeal an AI-driven decision, contact information for human review, data rights (access, correction, deletion), system limitations and appropriate reliance guidance. For B2B clients, includes integration documentation and client-configurable transparency settings.	End users, affected individuals, B2B client organizations, customer support teams, regulators
6. Post-Deployment Monitoring & Maintenance	Monitoring Dashboard & Audit Log	Maintains continuous records including: real-time performance metrics and drift detection alerts, explanation quality monitoring (are explanations still accurate after model updates?), complete decision audit	AI operations team, compliance officers, AI Ethics Board, executive leadership, external regulators (upon request)

AI Development Stage	Required Documentation	Detailed Contents & Purpose	Audience
		trail for high-risk systems, user feedback logs and complaint tracking, incident reports and remediation actions, model retraining documentation and before/after comparisons, quarterly transparency review reports, and annual third-party audit findings and remediation plans.	

Framework Sources: This documentation framework integrates best practices from Google's Model Cards for Model Reporting, Microsoft's Datasheets for Datasets and Responsible AI Impact Assessments, IBM's AI FactSheets, the EU AI Act's documentation requirements for high-risk AI systems (Articles 11-13), and the NIST AI RMF's "GOVERN" and "MAP" functions.

Slide 6: Concluding Reflections – Difficulties & Questions

The following reflections represent my personal observations and questions that emerged from studying AI explainability and transparency during this workshop. These are not AI-generated conclusions but genuine areas where I found the material challenging or thought-provoking.

Difficulties I Encountered

- **The audience problem is harder than it seems:** One of my biggest challenges was understanding how to create explanations that work for radically different audiences. A physician using a clinical decision support tool needs different information than a patient who was diagnosed using that same tool, and both differ from what a regulator needs during an audit. Creating a single explainability approach that serves all these audiences simultaneously feels nearly impossible—it seems like companies need layered explanation systems, which adds significant development complexity and cost.

Open Questions I'm Still Considering

- **How do we measure genuine understanding?** Providing an explanation is not the same as achieving comprehension. How can companies verify that users actually understand the explanations they receive, rather than just having information available that they ignore or misinterpret? What metrics or testing approaches can validate real comprehension?
- **Will explainability scale to foundation models?** Current XAI methods (SHAP, LIME, feature importance) were designed for traditional ML models. As companies

- **The accuracy-interpretability trade-off creates real ethical dilemmas:** I struggled with situations where the most accurate model (e.g., a deep neural network) is also the least explainable. In healthcare, do we accept slightly lower diagnostic accuracy for full explainability? Or do we accept a "black box" that saves more lives? Neither answer feels entirely satisfying, and the material didn't provide a clear resolution—perhaps because there isn't one.
- **Defining "good enough" transparency is subjective:** Different companies, regulators, and researchers seem to have very different standards for what constitutes adequate transparency. Google's approach focuses on user-level explanations, Microsoft emphasizes documentation and process, and the EU AI Act sets prescriptive requirements. Determining which standard NovaMind should follow—or how to reconcile them—was a persistent challenge throughout this assignment.

increasingly deploy large language models and multimodal foundation models with billions of parameters, will these techniques remain meaningful? Or do we need entirely new paradigms for explaining generative AI behavior?

- **How do you balance transparency with IP protection?** Full model transparency could expose proprietary architectures, training data compositions, or competitive advantages. How should companies like NovaMind navigate requests for transparency that might compromise trade secrets? Is there a middle ground between "full glass box" and "trust us"?
- **Who bears liability when explanations mislead?** If an AI system provides an explanation that a user relies on, but that explanation is oversimplified or technically inaccurate (a known limitation of post-hoc XAI methods), who is responsible for resulting harms? This legal question seems largely unresolved.

- **Retroactive transparency is problematic:** Much of the literature assumes transparency is designed in from the start, but many real-world AI systems were built without explainability in mind. Adding explanations to existing black-box models (via post-hoc methods like LIME or SHAP) raises questions about whether those explanations truly reflect the model's reasoning or just provide a plausible-sounding narrative.

- **Is "explainable AI" always desirable?** In some contexts (e.g., security systems, anti-fraud models), providing detailed explanations of how the AI works could enable adversaries to game the system. How should the RAI framework handle cases where transparency creates security vulnerabilities?

These difficulties and questions reinforce my belief that responsible AI is not a checkbox exercise but an ongoing process requiring continuous learning, stakeholder engagement, and willingness to make difficult trade-offs. The intersection of technical capability, ethical responsibility, regulatory compliance, and business reality makes this one of the most intellectually challenging—and important—aspects of working in AI today.